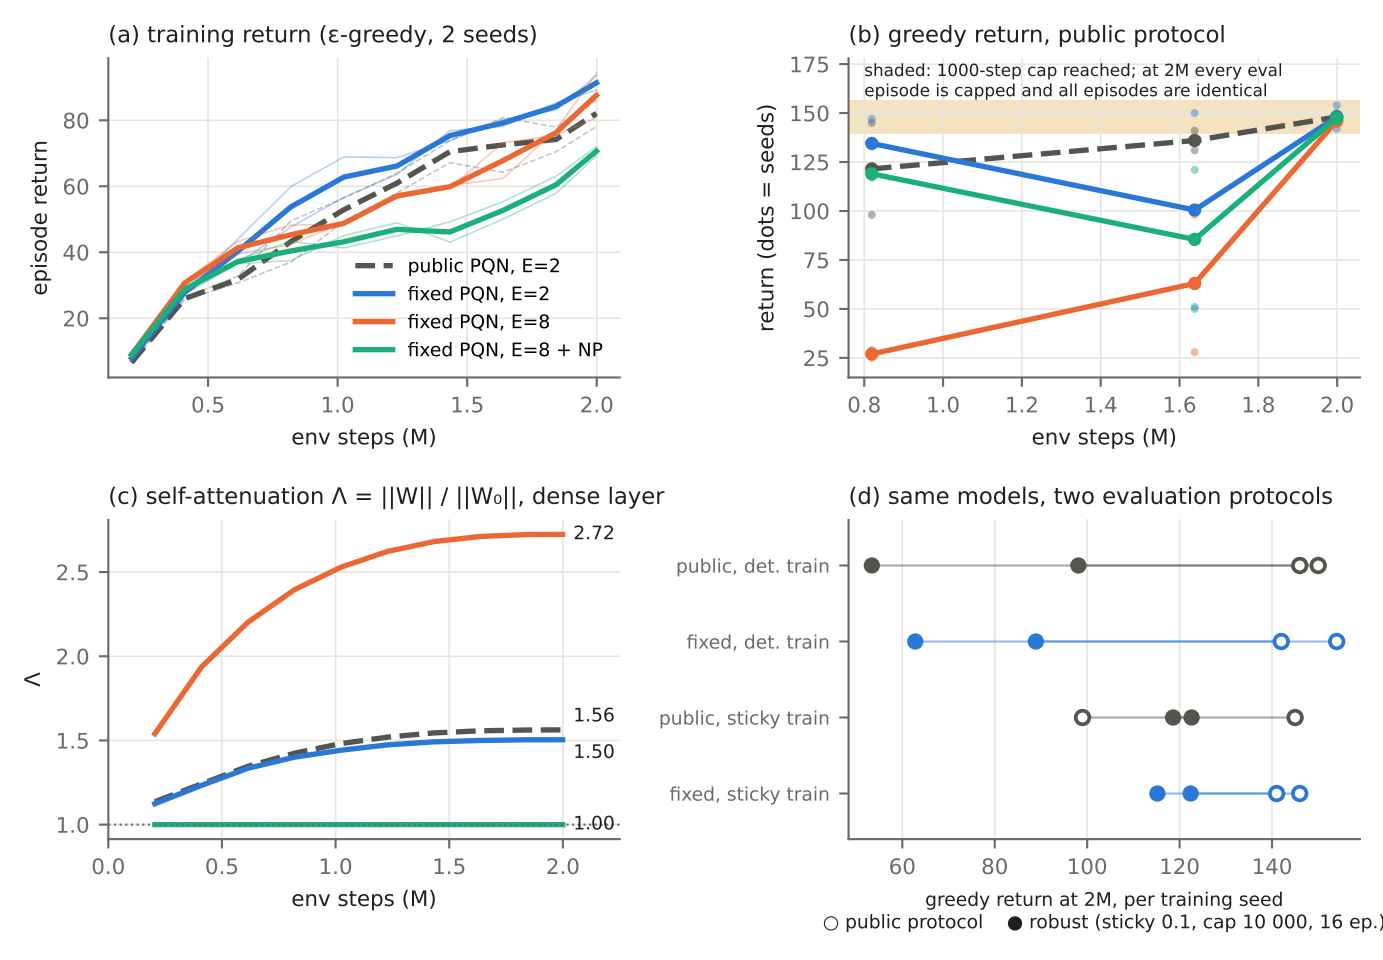
(a) training return (ε-greedy, 2 seeds)
episode return
env steps (M)
public PQN, E=2
fixed PQN, E=2
fixed PQN, E=8
fixed PQN, E=8 + NP
(b) greedy return, public protocol
shaded: 1000-step cap reached; at 2M every eval episode is capped and all episodes are identical
return (dots = seeds)
env steps (M)
(c) self-attenuation Λ = ||W|| / ||W₀||, dense layer
Λ
2.72
1.56
1.50
1.00
env steps (M)
(d) same models, two evaluation protocols
public, det. train
fixed, det. train
public, sticky train
fixed, sticky train
greedy return at 2M, per training seed
public protocol
robust (sticky 0.1, cap 10 000, 16 ep.)